

Assignment 2: Python for Analytics

- covers lectures 4-6
- due: November 8th by 6pm.
- Points will be deducted if:
 - Problems are not completed.
 - Portions of problems are not completed.
 - Third party modules were used when the question specified not to do so.
 - The problem was solved in a very inefficient manner. For instance, copying and pasting the same block of code 10 times instead of using a for loop or using a for loop when a comprehension would work.
 - Each day late will result in a 10% penalty.
 - Not attempting a problem or leaving it blank will result in 0 points for the problem and an additional 5 point deduction.

Question 1 (15 points)

Using the Iris data, sum the 4 numeric features and find out how many rows have a sum greater than 10. Do this in two ways.

- Using Numpy
- Using Pandas.

Print the shape for both the Pandas and Numpy solution.

```
In [1]: import sys
```

```
In [2]: print(sys.version)
```

```
3.9.12 (main, Apr  4 2022, 05:22:27) [MSC v.1916 64 bit (AMD64)]
```

```
In [1]: import numpy as np
import pandas as pd
import sklearn
from sklearn.datasets import load_iris
```

```
In [5]: print(np.__version__)
print(pd.__version__)
print(sklearn.__version__)
```

```
1.21.5
1.3.5
1.0.2
```

```
In [6]: iris = load_iris()
iris_df = pd.DataFrame(data= np.c_[iris['data'], iris['target']],
                        columns= iris['feature_names'] + ['target'])
iris_df['species'] = pd.Categorical.from_codes(iris.target, iris.target_names)
iris_df = iris_df.drop('target', 1)
```

C:\Users\Primo\AppData\Local\Temp\ipykernel_32412\1970229344.py:5: FutureWarning: In a future version of pandas all arguments of DataFrame.drop except for the argument 'labels' will be keyword-only
iris_df = iris_df.drop('target', 1)

```
In [7]: iris_df.head()
```

Out[7]:

	sepal length (cm)	sepal width (cm)	petal length (cm)	petal width (cm)	species
0	5.1	3.5	1.4	0.2	setosa
1	4.9	3.0	1.4	0.2	setosa
2	4.7	3.2	1.3	0.2	setosa
3	4.6	3.1	1.5	0.2	setosa
4	5.0	3.6	1.4	0.2	setosa

Use numpy

```
In [8]: cols = iris_df.columns
cols
```

Out[8]: Index(['sepal length (cm)', 'sepal width (cm)', 'petal length (cm)',
 'petal width (cm)', 'species'],
 dtype='object')

```
In [9]: np_sepal_length = iris_df['sepal length (cm)'].values
np_sepal_width = iris_df['sepal width (cm)'].values
np_petal_length = iris_df['petal length (cm)'].values
np_petal_width = iris_df['petal width (cm)'].values
sumOfStats = np_sepal_length + np_sepal_width + np_petal_length + np_petal_width
np_greater_than_10 = sumOfStats[sumOfStats > 10]
print(np_greater_than_10)
print("The shape is: ", np_greater_than_10.shape)
```

```
[10.2 10.2 11.4 10.1 10.8 11.2 12.  11.  10.3 11.5 10.7 10.7 10.7 10.6
 10.3 10.4 10.4 10.2 10.7 10.9 11.3 10.5 10.2 10.1 10.7 11.2 10.7 10.7
 16.3 15.6 16.4 13.1 15.4 14.3 15.9 11.6 15.4 13.2 11.5 14.6 13.2 15.1
 13.4 15.6 14.6 13.6 14.4 13.1 15.7 14.2 15.2 14.8 14.9 15.4 15.8 16.4
 14.9 12.8 12.8 12.6 13.6 15.4 14.4 15.5 16.  14.3 14.  13.3 13.7 15.1
 13.6 11.6 13.8 14.1 14.1 14.7 11.7 13.9 18.1 15.5 18.1 16.6 17.5 19.3
 13.6 18.3 16.8 19.4 16.8 16.3 17.4 15.2 16.1 17.2 16.8 20.4 19.5 14.7
 18.1 15.3 19.2 15.7 17.8 18.2 15.6 15.8 16.9 17.6 18.2 20.1 17.  15.7
 15.7 19.1 17.7 16.8 15.6 17.5 17.8 17.4 15.5 18.2 18.2 17.2 15.7 16.7
 17.3 15.8]
The shape is: (128,)
```

Use Pandas

```
In [10]: iris_df['sum of the stats'] = iris_df[cols[0:4]].sum(axis=1)
pd_greater_than_10 = iris_df.loc[iris_df['sum of the stats'] > 10]
print("The shape is: ", pd_greater_than_10.shape)
```

The shape is: (128, 6)

Question 2 (10 points)

Consider the below two arrays. The first will be actual values (y) and the second predicted values (\hat{y}). Calculate the below:

- MAE: Mean Absolute Error
 - defined as the average absolute error.
- MSE: Mean Squared Error
 - defined as taking the difference between the two arrays, squaring the errors, summing and finding the mean.
- MAPE: Mean Absolute Percentage Error
 - defined as the mean percentage difference between the two arrays.

Solve each using one line of code, making use of `numpy` array elementwise operations.

Print out each metric.

```
In [11]: y = np.array([1,4,5,2,4,6,1])
yhat = np.array([5,2,3,4,5,6,1])
```

MAE

```
In [12]: y_mae = np.mean(np.abs(y - yhat))  
y_mae
```

```
Out[12]: 1.5714285714285714
```

MSE

```
In [13]: y_mse = np.square(y - yhat).mean()  
y_mse
```

```
Out[13]: 4.142857142857143
```

MAPE

```
In [14]: y_mape = np.mean(np.abs((y - yhat) / y)) * 100  
y_mape
```

```
Out[14]: 87.85714285714286
```

Question 3 (10 points)

Find the standard deviation and mean of `sepal_length` using `describe()` and `loc`.

Use the above mean and standard deviation to create two variables:

- `upper_bound`, defined as `mean + 2 standard deviations`
- `lower_bound`, defined as `mean - 2 standard deviations`

Subset the dataframe for only rows where `sepal_length` is either greater than the `upper_bound` or less than the `lower_bound`.

Print the first 5 rows and shape of the subsetted dataframe.

```
In [16]: sepal_length_df = iris_df['sepal length (cm)']
sepal_length_mean = sepal_length_df.describe().loc['mean']
sepal_length_std = sepal_length_df.describe().loc['std']
upper = sepal_length_mean + 2 * sepal_length_std
lower = sepal_length_mean - 2 * sepal_length_std
subset_df = iris_df[iris_df['sepal length (cm)'].apply(lambda x: x > upper or x < lower)]
subset_df.head(5)
```

Out[16]:

	sepal length (cm)	sepal width (cm)	petal length (cm)	petal width (cm)	species	sum of the stats
105	7.6	3.0	6.6	2.1	virginica	19.3
117	7.7	3.8	6.7	2.2	virginica	20.4
118	7.7	2.6	6.9	2.3	virginica	19.5
122	7.7	2.8	6.7	2.0	virginica	19.2
131	7.9	3.8	6.4	2.0	virginica	20.1

Question 4 (15 points)

Load Boston Housing dataset from `sklearn` and put the data into a `pandas DataFrame` using the `data` and `feature_names` attributes from the `boston_data` object.

Find the IQR (interquartile range) for `AGE`, which is defined as the 75th quartile - the 25th quartile.

Remove observations with an `AGE` that are not within 1.5 IQR of the median. This means you will have to subset the data for less than `median + 1.5 IQR` and greater than `median - 1.5 IQR`.

Using the subsetted dataframe, find the strongest correlated feature with `AGE`, not including itself, and plot the two features as a scatter plot. Note strongest correlated could mean positive or negative.

Hint, this can be solved using the `corr()` method, finding the absolute value of the `corr` metric, and sorting.

Print the IQR, the highest correlating feature, the correlation itself and the scatter plot.

```
In [61]: from sklearn.datasets import load_boston
```

```
In [62]: boston_data = load_boston()
```

```
D:\Anaconda\lib\site-packages\sklearn\utils\deprecation.py:87: FutureWarning: Function load_boston is deprecated; `load_boston` is deprecated in 1.0 and will be removed in 1.2.
```

The Boston housing prices dataset has an ethical problem. You can refer to the documentation of this function for further details.

The scikit-learn maintainers therefore strongly discourage the use of this dataset unless the purpose of the code is to study and educate about ethical issues in data science and machine learning.

In this special case, you can fetch the dataset from the original source::

```
import pandas as pd
import numpy as np
```

```
data_url = "http://lib.stat.cmu.edu/datasets/boston"
raw_df = pd.read_csv(data_url, sep="\s+", skiprows=22, header=None)
data = np.hstack([raw_df.values[::2, :], raw_df.values[1::2, :2]])
target = raw_df.values[1::2, 2]
```

Alternative datasets include the California housing dataset (i.e. :func:`~sklearn.datasets.fetch_california_housing`) and the Ames housing dataset. You can load the datasets as follows::

```
from sklearn.datasets import fetch_california_housing
housing = fetch_california_housing()
```

for the California housing dataset and::

```
from sklearn.datasets import fetch_openml
housing = fetch_openml(name="house_prices", as_frame=True)
```

for the Ames housing dataset.

```
warnings.warn(msg, category=FutureWarning)
```

```
In [169]: df = pd.DataFrame(boston_data['data'], columns=boston_data['feature_names'])
df
```

Out[169]:

	CRIM	ZN	INDUS	CHAS	NOX	RM	AGE	DIS	RAD	TAX	PTRATIO	B	LST
0	0.00632	18.0	2.31	0.0	0.538	6.575	65.2	4.0900	1.0	296.0	15.3	396.90	4.
1	0.02731	0.0	7.07	0.0	0.469	6.421	78.9	4.9671	2.0	242.0	17.8	396.90	9.
2	0.02729	0.0	7.07	0.0	0.469	7.185	61.1	4.9671	2.0	242.0	17.8	392.83	4.
3	0.03237	0.0	2.18	0.0	0.458	6.998	45.8	6.0622	3.0	222.0	18.7	394.63	2.
4	0.06905	0.0	2.18	0.0	0.458	7.147	54.2	6.0622	3.0	222.0	18.7	396.90	5.
...
501	0.06263	0.0	11.93	0.0	0.573	6.593	69.1	2.4786	1.0	273.0	21.0	391.99	9.
502	0.04527	0.0	11.93	0.0	0.573	6.120	76.7	2.2875	1.0	273.0	21.0	396.90	9.
503	0.06076	0.0	11.93	0.0	0.573	6.976	91.0	2.1675	1.0	273.0	21.0	396.90	5.
504	0.10959	0.0	11.93	0.0	0.573	6.794	89.3	2.3889	1.0	273.0	21.0	393.45	6.
505	0.04741	0.0	11.93	0.0	0.573	6.030	80.8	2.5050	1.0	273.0	21.0	396.90	7.

506 rows × 13 columns

```
In [173]: df_desc = df['AGE'].describe()
print("The 25% quantile is: ", df_desc.loc['25%'])
print("The 50% quantile is: ", df_desc.loc['50%'])
print("The 75% quantile is: ", df_desc.loc['75%'])
IQR = df_desc.loc['75%'] - df_desc.loc['25%']
print("IQR is: ", IQR)
```

```
The 25% quantile is: 45.025
The 50% quantile is: 77.5
The 75% quantile is: 94.07499999999999
IQR is: 49.04999999999999
```

```
In [176]: median_upper = df_desc.loc['50%'] + 1.5 * IQR
median_lower = df_desc.loc['50%'] - 1.5 * IQR

df = df[df['AGE'].apply(lambda x: x > median_lower and x < median_upper)]
df
```

Out[176]:

	CRIM	ZN	INDUS	CHAS	NOX	RM	AGE	DIS	RAD	TAX	PTRATIO	B	LST
0	0.00632	18.0	2.31	0.0	0.538	6.575	65.2	4.0900	1.0	296.0	15.3	396.90	4.
1	0.02731	0.0	7.07	0.0	0.469	6.421	78.9	4.9671	2.0	242.0	17.8	396.90	9.
2	0.02729	0.0	7.07	0.0	0.469	7.185	61.1	4.9671	2.0	242.0	17.8	392.83	4.
3	0.03237	0.0	2.18	0.0	0.458	6.998	45.8	6.0622	3.0	222.0	18.7	394.63	2.
4	0.06905	0.0	2.18	0.0	0.458	7.147	54.2	6.0622	3.0	222.0	18.7	396.90	5.
...
501	0.06263	0.0	11.93	0.0	0.573	6.593	69.1	2.4786	1.0	273.0	21.0	391.99	9.
502	0.04527	0.0	11.93	0.0	0.573	6.120	76.7	2.2875	1.0	273.0	21.0	396.90	9.
503	0.06076	0.0	11.93	0.0	0.573	6.976	91.0	2.1675	1.0	273.0	21.0	396.90	5.
504	0.10959	0.0	11.93	0.0	0.573	6.794	89.3	2.3889	1.0	273.0	21.0	393.45	6.
505	0.04741	0.0	11.93	0.0	0.573	6.030	80.8	2.5050	1.0	273.0	21.0	396.90	7.

505 rows × 13 columns




```
In [180]: df_corr = abs(df.corr())
df_corr
```

Out[180]:

	CRIM	ZN	INDUS	CHAS	NOX	RM	AGE	DIS	RAI
CRIM	1.000000	0.200941	0.406307	0.056124	0.420654	0.218830	0.352826	0.379311	0.625355
ZN	0.200941	1.000000	0.534751	0.042973	0.518052	0.312880	0.575041	0.666002	0.312924
INDUS	0.406307	0.534751	1.000000	0.062633	0.763456	0.391164	0.645664	0.707763	0.594768
CHAS	0.056124	0.042973	0.062633	1.000000	0.090788	0.091675	0.085728	0.098771	0.007780
NOX	0.420654	0.518052	0.763456	0.090788	1.000000	0.301322	0.731783	0.768844	0.610923
RM	0.218830	0.312880	0.391164	0.091675	0.301322	1.000000	0.238471	0.204259	0.209032
AGE	0.352826	0.575041	0.645664	0.085728	0.731783	0.238471	1.000000	0.748317	0.455260
DIS	0.379311	0.666002	0.707763	0.098771	0.768844	0.204259	0.748317	1.000000	0.493910
RAD	0.625353	0.312924	0.594768	0.007780	0.610923	0.209032	0.455260	0.493910	1.000000
TAX	0.582648	0.315982	0.720533	0.036190	0.667401	0.291073	0.504915	0.533560	0.910166
PTRATIO	0.289805	0.392045	0.383103	0.121671	0.188635	0.355341	0.261762	0.232213	0.464641
B	0.384911	0.175883	0.356761	0.048967	0.379833	0.127710	0.273585	0.291211	0.444231
LSTAT	0.455357	0.414643	0.603406	0.054589	0.590077	0.613328	0.601240	0.496011	0.487891


```
In [186]: a = np.argsort(df_corr['AGE'].values)
index = df_corr.columns
highest_corr = index[a[-2]]
highest_corr
```

Out[186]: 'DIS'

```
In [191]: import matplotlib.pyplot as plt
import seaborn as sns
x = df['AGE'].values
y = df['DIS'].values
plt.figure(figsize=(10,5))
ax = sns.scatterplot(x,y)
ax.set_xlabel("AGE", fontsize = 10)
ax.set_ylabel("DIS", fontsize = 10)
ax.set_title("AGE vs. DIS", fontsize = 15)
plt.show()
```

D:\Anaconda\lib\site-packages\seaborn_decorators.py:36: FutureWarning: Pass the following variables as keyword args: x, y. From version 0.12, the only valid positional argument will be `data`, and passing other arguments without an explicit keyword will result in an error or misinterpretation.

warnings.warn(



Question 5 (10 points)

rating_df is a rating matrix, where each row is a user, each column is a movie, and the cells are how a given user rated a given movie. For instance, the upper left cell has a 4, meaning user_1 rated star_wars as a 4.

Perform the below 3 transformations on rating_df.

- min_max: 0-1 scale
 - defined as $(x - \min(x)) / (\max(x) - \min(x))$
- mean_centered:
 - $x - \text{mean}(x)$
- z_score:
 - $(x - \text{mean}(x)) / \text{std}(x)$

This means, for instance, each column for min max should be scaled to where the max is 1 and the min is 0.

Hint, this should be done using 1 line, making use of broadcasting and rows and columnwise mean, min, max and standard deviation calculations.

DO NOT USE ANY FUNCTIONALITY IMPORTED FROM `sklearn`.

Print out the 3 scaled dataframes.

```
In [155]: import numpy as np
import pandas as pd

user_1 = np.array([4,2,5])
user_2 = np.array([1,5,4])
user_3 = np.array([2,4,2])
user_4 = np.array([3,5,4])

rating_matrix = np.array([user_1, user_2, user_3, user_4])

columns = ["star_wars", "harry_potter", "avengers"]
index = ["user_1", "user_2", "user_3", "user_4"]

rating_df = pd.DataFrame(rating_matrix, columns = columns, index = index)

rating_df
```

Out[155]:

	star_wars	harry_potter	avengers
user_1	4	2	5
user_2	1	5	4
user_3	2	4	2
user_4	3	5	4

```
In [161]: rating_df_min_max = (rating_df[['star_wars', 'harry_potter', 'avengers']] - \
                                rating_df[['star_wars', 'harry_potter', 'avengers']].min())
                                (rating_df[['star_wars', 'harry_potter', 'avengers']].max() -
                                rating_df[['star_wars', 'harry_potter', 'avengers']].min())
rating_df_min_max
```

Out[161]:

	star_wars	harry_potter	avengers
user_1	1.000000	0.000000	1.000000
user_2	0.000000	1.000000	0.666667
user_3	0.333333	0.666667	0.000000
user_4	0.666667	1.000000	0.666667

```
In [164]: rating_df_mean_centered = rating_df[['star_wars', 'harry_potter', 'avengers']] - \
        rating_df[['star_wars', 'harry_potter', 'avengers']].mean()
rating_df_mean_centered
```

Out[164]:

	star_wars	harry_potter	avengers
user_1	1.5	-2.0	1.25
user_2	-1.5	1.0	0.25
user_3	-0.5	0.0	-1.75
user_4	0.5	1.0	0.25

```
In [166]: rating_df_z_score = (rating_df[['star_wars', 'harry_potter', 'avengers']] - \
        rating_df[['star_wars', 'harry_potter', 'avengers']].mean()) / \
        rating_df[['star_wars', 'harry_potter', 'avengers']].std()
rating_df_z_score
```

Out[166]:

	star_wars	harry_potter	avengers
user_1	1.161895	-1.414214	0.993399
user_2	-1.161895	0.707107	0.198680
user_3	-0.387298	0.000000	-1.390759
user_4	0.387298	0.707107	0.198680

Quesiton 6 (15 points)

Find the pariwise distances of each users rating vector using the eudclidean distance. For instance, user_1 has a vector of [4,2,5] while user_2 has a vector of [1,5,4] . Finding the distance between these two vectors would give us the distance between user_1 and user_2 .

Add a column to rating_df called most_similar_user that has the user_id of the most similar user for that given observation.

Note, when making a distance matrix, the min distance is going to be the distance between each user and themselves. Make sure the most_similar_user is not the user themself.

Hint, this can be solved using squareform and pdist from scipy , then sorting the resulting distance matrix using argsort from numpy . Remember, argsort sorts the values, then provides an index, so the index can then be converted to a user using the columns list of users.

Print out the dataframe with the new column.

```
In [148]: import pandas as pd
import numpy as np
from scipy.spatial.distance import pdist, squareform
```

```
In [149]: user_1 = np.array([4,2,5])
user_2 = np.array([1,5,4])
user_3 = np.array([2,4,2])
user_4 = np.array([3,5,4])

rating_matrix = np.array([user_1, user_2, user_3, user_4])

columns = ["star_wars", "harry_potter", "avengers"]
index = ["user_1", "user_2", "user_3", "user_4"]

rating_df = pd.DataFrame(rating_matrix, columns = columns, index = index)

rating_df
```

Out[149]:

	star_wars	harry_potter	avengers
user_1	4	2	5
user_2	1	5	4
user_3	2	4	2
user_4	3	5	4

```
In [150]: dist_df = pd.DataFrame(squareform(pdist(rating_df)), index = index, columns = index)

dist_df
```

Out[150]:

	user_1	user_2	user_3	user_4
user_1	0.000000	4.358899	4.123106	3.316625
user_2	4.358899	0.000000	2.449490	2.000000
user_3	4.123106	2.449490	0.000000	2.449490
user_4	3.316625	2.000000	2.449490	0.000000

```
In [154]: lowest_dist = []
for row in index:
    a = np.argsort(dist_df[row].values)
    lowest_dist.append(index[a[1]])

rating_df['most similarity'] = lowest_dist
rating_df
```

Out[154]:

	star_wars	harry_potter	avengers	most similarity
user_1	4	2	5	user_4
user_2	1	5	4	user_4
user_3	2	4	2	user_2
user_4	3	5	4	user_2

Question 7 (10 points)

Use a for loop to make a 2,3 and 4 period rolling mean column for each user. Making sure to add each column to the dataframe.

Hint, since this is finding the rolling mean for each user, we can use `groupby` and `rolling` in `pandas`.

Print the dataframe out.

```
In [129]: import numpy as np
import pandas as pd
```

```
In [130]: def my_func(x):
            if x == 1:
                return x
```

```
In [131]: [my_func(i) for i in [1,2,3,4,5]]
```

```
Out[131]: [1, None, None, None, None]
```

```
In [132]: metric = np.array([5,3,2,4,5,1,4,1,4,2,5,3,1,2,3])
ids = np.array(["a","a","a","a","a","b","b","b","b","b","c","c","c","c","c"])

df = pd.DataFrame({
    "id":ids,
    "metric":metric
})
df
```

```
Out[132]:
```

	id	metric
0	a	5
1	a	3
2	a	2
3	a	4
4	a	5
5	b	1
6	b	4
7	b	1
8	b	4
9	b	2
10	c	5
11	c	3
12	c	1
13	c	2
14	c	3

```
In [134]: rolling_metrics = [2,3,4]
for i in rolling_metrics:
    name = "rolling_metric_" + str(i)
    a = df.groupby(['id']).rolling(i).mean().reset_index()
    df[name] = a['metric']

df
```

Out[134]:

	id	metric	rolling_metric_2	rolling_metric_3	rolling_metric_4
0	a	5	NaN	NaN	NaN
1	a	3	4.0	NaN	NaN
2	a	2	2.5	3.333333	NaN
3	a	4	3.0	3.000000	3.50
4	a	5	4.5	3.666667	3.50
5	b	1	NaN	NaN	NaN
6	b	4	2.5	NaN	NaN
7	b	1	2.5	2.000000	NaN
8	b	4	2.5	3.000000	2.50
9	b	2	3.0	2.333333	2.75
10	c	5	NaN	NaN	NaN
11	c	3	4.0	NaN	NaN
12	c	1	2.0	3.000000	NaN
13	c	2	1.5	2.000000	2.75
14	c	3	2.5	2.000000	2.25

Question 8 (15 points)

The below dataframe has the `sales` in each `month` for 3 `products`. The first 5 rows of data can be interpreted as monthly sales for months 1-5 for product `a`. So, the `sales` for product `a` are `[5,3,2,4,5]`.

Pivot the below dataframe so the rows are the `month`, the columns are the `products` and the cell values are the `sales` for a given `month-product`.

Find pairwise correlations for each `products` `sales`. The result should be a 3 x 3 correlation matrix.

Find the `products` with the highest correlating sales. Create a dataframe with two columns, the first being an `product` and the second column the highest correlating product.

Hint, once the data is pivoted, you can use the `corr()` method, get the absolute value of the correlation and use `argsort` to sort the correlation dataframe. Remember, highest correlation could be positive or negative.

Print the 3 by 3 correlation matrix and the two column dataframe with the most similar ids.

```
In [18]: metric = np.array([5,3,2,4,5,1,4,1,4,2,5,3,1,2,3])
periods = [1,2,3,4,5] * 3
ids = np.array(["a","a","a","a","a","b","b","b","b","b","c","c","c","c","c"])

df = pd.DataFrame({
    "product": ids,
    "sales": metric,
    "month": periods
})

df
```

Out[18]:

	product	sales	month
0	a	5	1
1	a	3	2
2	a	2	3
3	a	4	4
4	a	5	5
5	b	1	1
6	b	4	2
7	b	1	3
8	b	4	4
9	b	2	5
10	c	5	1
11	c	3	2
12	c	1	3
13	c	2	4
14	c	3	5

```
In [19]: pivot_df = df.pivot(index = "month", columns = "product", values = "sales")
pivot_df
```

Out[19]:

product	a	b	c
month			
1	5	1	5
2	3	4	3
3	2	1	1
4	4	4	2
5	5	2	3


```
In [20]: df_corr = abs(pivot_df.corr())  
df_corr
```

Out[20]:

	product	a	b	c
product				
a	1.000000	0.075858	0.749777	
b	0.075858	1.000000	0.177822	
c	0.749777	0.177822	1.000000	

```
In [21]: index = df_corr.index  
highest_corr = []  
for row in index:  
    a = np.argsort(df_corr[row].values)  
    highest_corr.append(index[a[-2]])  
  
highest_corr_d = {'col1':index, 'highest correlation product':highest_corr}  
highest_corr_df = pd.DataFrame(highest_corr_d)  
highest_corr_df
```

Out[21]:

	col1	highest correlation product
0	a	c
1	b	c
2	c	a